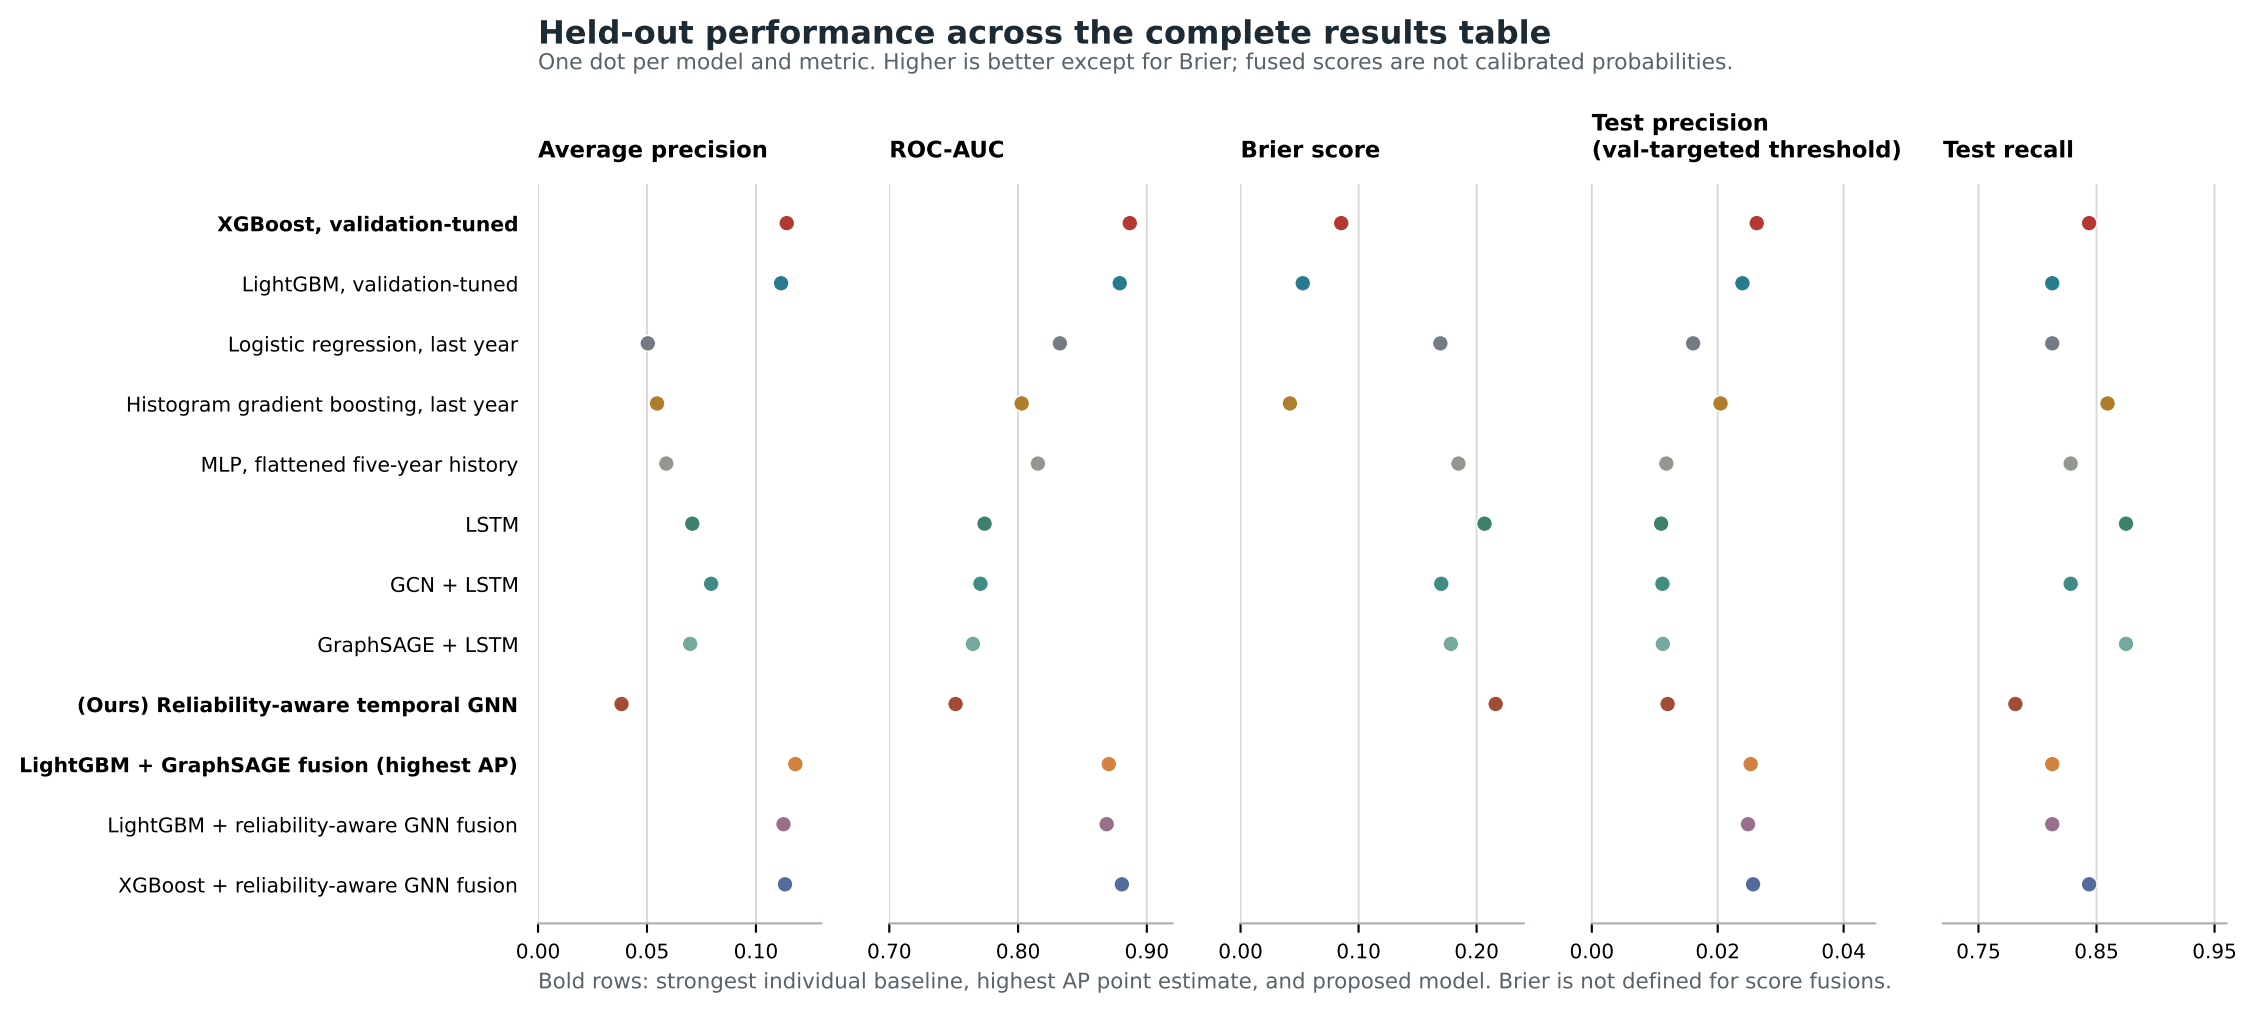
Held-out performance across the complete results table
One dot per model and metric. Higher is better except for Brier; fused scores are not calibrated probabilities.
Average precision
ROC-AUC
Brier score
Test precision (val-targeted threshold)
Test recall
XGBoost, validation-tuned
LightGBM, validation-tuned
Logistic regression, last year
Histogram gradient boosting, last year
MLP, flattened five-year history
LSTM
GCN + LSTM
GraphSAGE + LSTM
(Ours) Reliability-aware temporal GNN
LightGBM + GraphSAGE fusion (highest AP)
LightGBM + reliability-aware GNN fusion
XGBoost + reliability-aware GNN fusion
0.00
0.05
0.10
0.70
0.80
0.90
0.00
0.10
0.20
0.00
0.02
0.04
0.75
0.85
0.95
Bold rows: strongest individual baseline, highest AP point estimate, and proposed model. Brier is not defined for score fusions.